

Bayesian Statistics

Generalized Linear Mixed Effects Model

Nan Lin

Department of Mathematics


Washington University in St. Louis

Generalized Linear Mixed Effects Model (GLMM)

- ▶ Extension of Linear mixed effects model to non-normal response
 - ▶ e.g. binary or count

$$\beta_1, \dots, \beta_m \sim \text{i.i.d. multivariate normal}(\theta, \Sigma)$$

$$p(\mathbf{y}_j | \mathbf{X}_j, \beta_j, \gamma) = \prod_{i=1}^{n_j} p(y_{i,j} | \beta_j^T \mathbf{x}_{i,j}, \gamma),$$



Metropolis-within-Gibbs

- ▶ For simplicity, we assume no γ parameter (such as in logistic and Poisson regression)
- ▶ Full conditional distributions of (θ, Σ) can be standard under proper prior specifications

- ▶ Prior

$$\theta \sim \text{multivariate normal}(\mu_0, \Lambda_0)$$

$$\Sigma \sim \text{inverse-Wishart}(\eta_0, \mathbf{S}_0^{-1})$$

- ▶ Full conditional distributions are of the same type
- ▶ $\{\beta_j\}$ requires Metropolis updating
 - ▶ We consider updating $\{\beta_j\}$ sequentially one at a time

Metropolis updating of $\{\beta_j\}$

- ▶ Consider a multivariate normal proposal distribution with mean as the current value $\beta_j^{(s)}$ with proposal variance $V_j^{(s)}$

- ▶ A common choice is to set $V_j^{(s)} = \text{scale} \times \Sigma^{(s)}$

1. Sample $\beta_j^* \sim \text{multivariate normal}(\beta_j^{(s)}, V_j^{(s)})$.
2. Compute the acceptance ratio

$$r = \frac{p(\mathbf{y}_j | \mathbf{X}_j, \beta_j^*) p(\beta_j^* | \boldsymbol{\theta}^{(s)}, \Sigma^{(s)})}{p(\mathbf{y}_j | \mathbf{X}_j, \beta_j^{(s)}) p(\beta_j^{(s)} | \boldsymbol{\theta}^{(s)}, \Sigma^{(s)})}.$$

3. Sample $u \sim \text{uniform}(0,1)$. Set $\beta_j^{(s+1)}$ to β_j^* if $u < r$ and to $\beta_j^{(s)}$ if $u > r$.

Putting together

1. Sample $\boldsymbol{\theta}^{(s+1)}$ from its full conditional distribution.
2. Sample $\Sigma^{(s+1)}$ from its full conditional distribution.
3. For each $j \in \{1, \dots, m\}$,
 - a) propose a new value β_j^* ;
 - b) set $\beta_j^{(s+1)}$ equal to β_j^* or $\beta_j^{(s)}$ with the appropriate probability.

Example: tumor location data*

- ▶ A certain population of laboratory mice experiences a high rate of intestinal tumor growth. One item of interest to researchers is how the rate of tumor growth varies along the length of the intestine.
- ▶ Study: the intestine of each of 21 sample mice was divided into 20 sections and the number of tumors occurring in each section was recorded.
- ▶ Code the 20 locations as $x \in (0.05, 0.10, \dots, 0.95)$
- ▶ Response: # of tumors

*From Haigis KM, Hoff PD, White A, Shoemaker AR, Halberg RB, Dove VF (2004) Tumor regionality in the mouse intestine reflects the mechanism of loss of apc function. *PNAS* 101(26):9769–9773

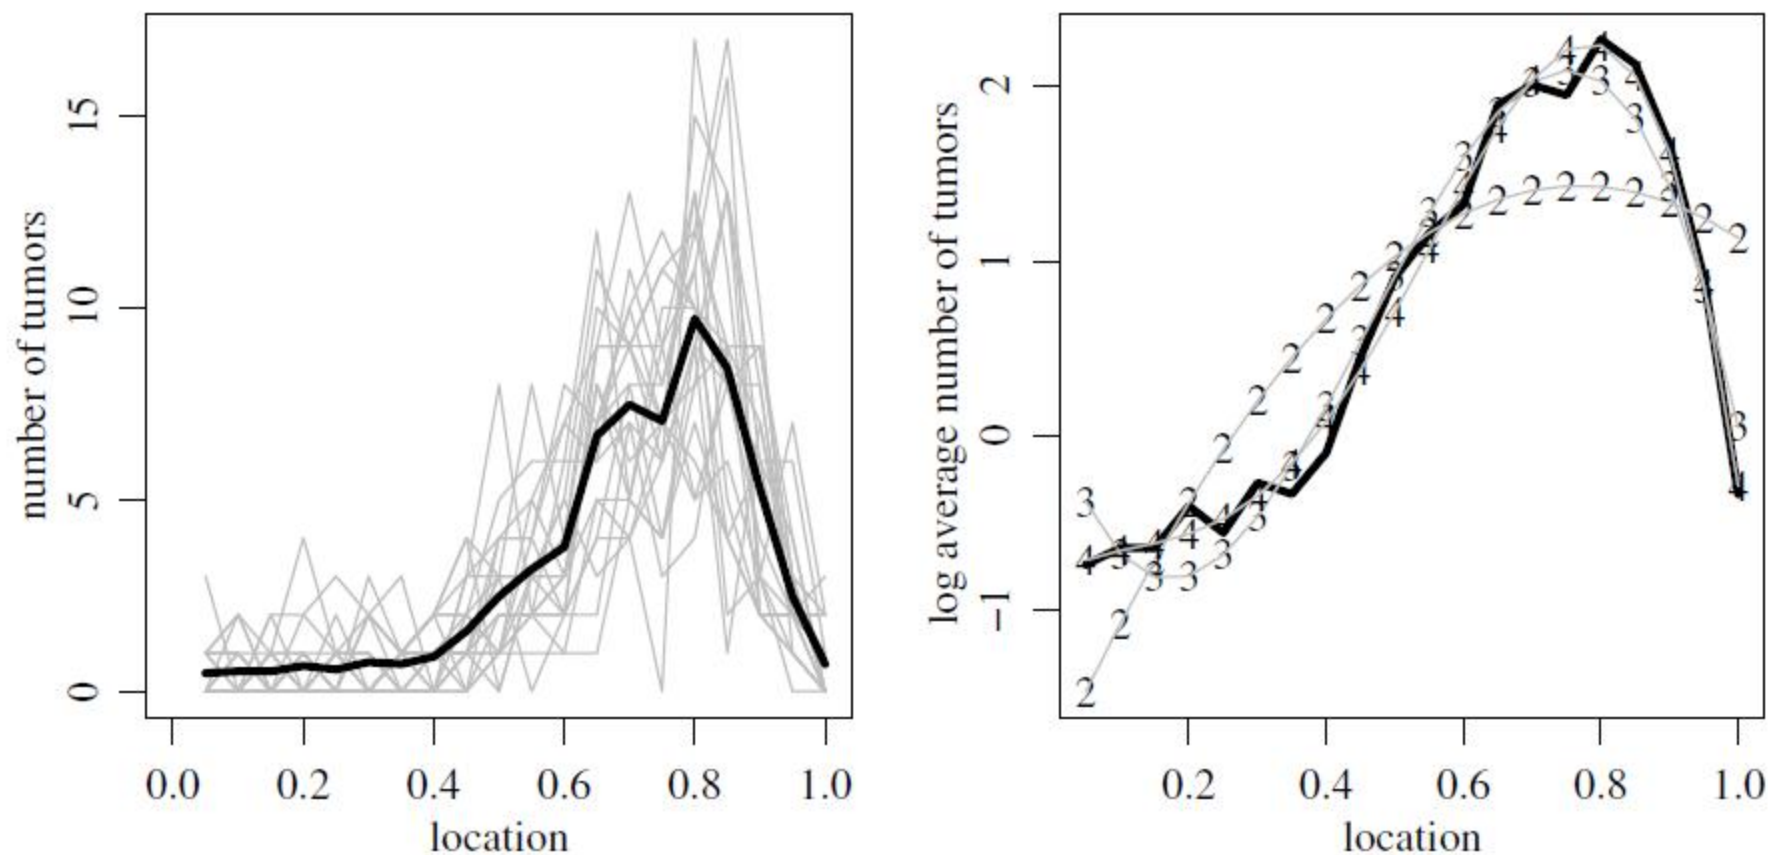


Fig. 11.4. Tumor count data. The first panel gives mouse-specific tumor counts as a function of location in gray, with a population average in black. The second panel gives quadratic, cubic and quartic polynomial fits to the log sample average tumor count.

Hierarchical Poisson regression

- ▶ Let $Y_{x,j}$ be mouse j 's tumor count at location x of their intestine
- ▶ Level 1:
 - ▶ Assume $Y_{x,j} \sim \text{Poisson}(e^{f_j(x)})$
 - ▶ $f_j(x) = \beta_{1,j} + \beta_{2,j}x + \cdots + \beta_{p,j}x^{p-1}$, use $p = 5$

- ▶ Level 2:

$$\beta_1, \dots, \beta_m \sim \text{i.i.d. multivariate normal}(\boldsymbol{\theta}, \Sigma)$$

- ▶ Level 3:

$$\boldsymbol{\theta} \sim \text{multivariate normal}(\boldsymbol{\mu}_0, \Lambda_0)$$

$$\Sigma \sim \text{inverse-Wishart}(\eta_0, \mathbf{S}_0^{-1})$$

Choosing the hyper-prior

- ▶ Unit information prior
 - ▶ the prior distributions for the parameters are weakly centered around estimates derived from the observed data
- ▶ Get a rough estimate $\tilde{\beta}_j$ by fitting a polynomial regression using data from mouse j only
 - ▶ Regress $\left\{ \log \left(y_{1,j} + \frac{1}{n} \right), \dots, \log \left(y_{n,j} + \frac{1}{n} \right) \right\}$ on $\{x_1, \dots, x_{20}\}$
- ▶ Set $\mu_0 = \frac{1}{m} \sum_j \tilde{\beta}_j$ and Λ_0 as the sample covariance matrix from $\tilde{\beta}_j$'s
- ▶ Set S_0 also as the sample covariance matrix from $\tilde{\beta}_j$'s
- ▶ $\eta_0 = p + 2 = 7$

MCMC

- ▶ Proposal distribution for β_j : multivariate normal $(\beta_j^{(s)}, \Sigma^{(s)} / 2)$
 - ▶ Acceptance rate is around 31%
- ▶ 50,000 scans
- ▶ Thinning: every 10th scan

```

## data
data( chapter11 )
Y<-XY.tumor$Y ; X<-XY.tumor$X ; m<-dim(Y)[1] ; p<-dim(X)[2]

## priors
BETA<-NULL
for( j in 1:m)
{
  BETA<-rbind( BETA, lm( log( Y[j,]+1/20) ~ -1+X[, , j] ) $coef)
}

mu0<-apply( BETA, 2 , mean)
S0<-cov( BETA) ; eta0<-p+2
iL0<-iSigma<-solve( S0)

## MCMC
THETA.post<-NULL ; set.seed(1)
for( s in 1:50000)
{
  ##update theta
  Lm<-solve( iL0 + m*iSigma )
  mum<-Lm%*%( iL0%*%mu0 + iSigma%*%apply( BETA, 2 , sum) )
  theta<-t( rmvnorm( 1 , mum, Lm) )
  ##

  ##update Sigma
  mtheta<-matrix( theta , m, p , byrow=TRUE)
  iSigma<-rwish( 1 , eta0+m,
                 solve( S0+t( BETA-mtheta)%*%( BETA-mtheta) ) )
  ##

```

```

##update beta
Sigma<-solve(iSigma) ; dSigma<-det(Sigma)
for(j in 1:m)
{
  beta.p<-t(rmvnorm(1,BETA[j,],.5*Sigma))

  lr<-sum( dpois(Y[j,],exp(X[,j]%*%beta.p),log=TRUE ) -
          dpois(Y[j,],exp(X[,j]%*%BETA[j,]),log=TRUE ) ) +
          ldmvnorm( t(beta.p),theta,Sigma,
                    iSigma=iSigma,dSigma=dSigma ) -
          ldmvnorm( t(BETA[j,]),theta,Sigma,
                    iSigma=iSigma,dSigma=dSigma )

  if( log(runif(1))<lr ) { BETA[j,]<-beta.p }
}
##

##store some output
if(s%%10==0){THETA.post<-rbind(THETA.post,t(theta))}
##
}

```

Results




Fig. 11.5. 2.5, 50 and 97.5% quantiles for $\exp(\theta^T x)$, $\exp(\beta^T x)$ and $\{Y|x\}$.

The widening confidence bands of the three plots in this figure describe cumulative sources of uncertainty

- Uncertainty in θ + across-mouse heterogeneity + fluctuation in the response around its mean

Discussion

- ▶ Posterior approximation via MCMC for hierarchical models can suffer from poor mixing.
 - ▶ Many of the parameters in the model are highly correlated, and generating them one at a time in the Gibbs sampler can lead to a high degree of autocorrelation
- ▶ **R package** `MCMCglmm`
 - ▶ [Tutorial of the package](#)