

Bayesian Statistics

Simulation

Nan Lin

Department of Mathematics

Washington University in St. Louis

Motivation

- ▶ Posterior distribution is often known up to an unknown normalizing constant
- ▶ How do we compute posterior mean, median or variance based on a not-fully-known distribution?
- ▶ Simulation is a general approach to solve the problem if we could generate random numbers from the posterior distribution.

Law of large numbers (LLN)

- ▶ Given a random sample of size n , sample mean converges in probability to the population mean as the sample size n approaches infinity.
- ▶ Application:
 - ▶ Suppose that we are interested in the random variable $Y = g(X)$ and X has distribution is $p(x)$
 - ▶ Treat it as the population distribution and generate a large set of independent random numbers from $p(x)$, x_1, \dots, x_B
 - ▶ Then $y_i = g(x_i)$ are a set of independent random numbers of from the distribution of Y .
 - ▶ For example, to learn $E(Y)$, we can approximate it by \bar{y} . Based on the LLN, it can be made very accurate by making B large.

Random number generation: discrete distributions

- ▶ Suppose the distribution is supported at $\{1, 2, \dots, k\}$ and the corresponding probabilities are p_i .
 - ▶ Note that the values in the support have no numerical meaning. They are just a set of labels.
- ▶ We can think the problem as randomly sampling some subjects from a population with k subjects.
 - ▶ But note, the same subject can appear multiple times, so, this is “sampling with replacement”.
 - ▶ Second, in basic sampling setup, every subject is given equal chance to be selected. It is NOT the case here.
 - ▶ So, our problem is “weighted random sampling with replacement”.

How to do sampling in R?

- ▶ Use the function `sample()`
- ▶ **Usage:** `sample(x, size, replace = FALSE, prob = NULL)`
- ▶ In our notation,
 - ▶ $x: \{1, 2, \dots, k\}$
 - ▶ size: B
 - ▶ `replace: TRUE`
 - ▶ `prob: $\{p_1, \dots, p_k\}$`
- ▶ Exercise

Random number generation: continuous distributions

- ▶ Probability integral transformation
 - ▶ If a continuous random variable X has cdf $F(x)$, then $U = F(X)$ has distribution $U(0,1)$.
- ▶ If we start from U , and F is invertible, we will have $X = F^{-1}(U)$.
- ▶ Almost all computer languages have a random number generator for $U(0,1)$.
- ▶ So, we just need to figure out F^{-1} .
- ▶ Exercise: simulate random numbers from an exponential distribution with mean $\lambda = 1$.

How to deal with the unknown normalizing constant?

- ▶ The approaches we discussed so far all require the distribution is fully known.
- ▶ But in Bayesian inference, we usually do not know the normalizing constant. What do we do then?
- ▶ Let's see the following example.


Discrete prior

► Model: $X \sim \text{Binomial}(n, p)$

► $x = 11, n - x = 16$


► Support of p :

► $\{0.05 \ 0.15 \ 0.25 \ 0.35 \ 0.45 \ 0.55 \ 0.65 \ 0.75 \ 0.85 \ 0.95\}$




Histogram prior

► Step function



Posterior



A “brute-force” method

- ▶ How to summarize posterior computation for an arbitrary prior density $g(p)$?
 - ▶ Choose a grid of values of p over an interval that covers the posterior density
 - ▶ Compute the product of the likelihood $L(p)$ and the prior $g(p)$ on the grid
 - ▶ Normalize by dividing each product by the sum of the products. In this step, the posterior density is approximated by a discrete probability distribution on the grid
 - ▶ Using the R function `sample()`, take a random sample with replacement from the discrete distribution.
 - ▶ The resulting simulated draws are an approximate sample from the posterior distribution.